

Introduction to Computer Networks

Fair Queuing (§5.4.3)




David Wetherall (djw@uw.edu)

Professor of Computer Science & Engineering

UNIVERSITY *of* WASHINGTON

Topic

- Sharing bandwidth between flows
- ➔ WFQ (Weighted Fair Queuing)
 - Key building block for QOS




Sharing with FIFO Queuing

- FIFO “drop tail” queue:
 - Queue packets First In First Out (FIFO)
 - ➔ Discard new packets when full
 - Typical router queuing model
- Sharing with FIFO queue
 - ➔ Multiple users or flows send packets over the same (output) link
 - What will happen?


Sharing with FIFO Queuing (2)

- Bandwidth allocation depends on behavior of all flows
- TCP gives long-term sharing – with delay/loss, and RTT bias
- Aggressive user/flow can crowd out the others




Round-Robin Queuing

- Idea to improve fairness:
 - Queue packets separately for each flow; take one packet in turn from each non-empty flow at the next output time




Round-Robin Queuing (2)

- Idea to improve fairness:
 - Queue packets separately for each flow; take one packet in turn from each non-empty flow at the next output time
 - How well does this work?



Round-Robin Queuing (3)

- Flows don't see uncontrolled delay/loss from others!
- But different packet sizes lead to bandwidth imbalance
 - Might be significant, e.g., 40 bytes vs 1500 bytes



➤ Fair Queuing

- Round-robin but approximate bit-level fairness:
 - Approximate by computing virtual finish time
 - Virtual clock ticks once for each bit sent from all flows
 - Send packets in order of their virtual finish times, $\text{Finish}(j)_F$
 - Not perfect – don't preempt packet being transmitted

➤ $\text{Arrive}(j)_F$ = arrival time of j-th packet of flow F


➤ $\text{Length}(j)_F$ = length of j-th packet of flow F

➤ $\text{Finish}(j)_F = \max(\text{Arrive}(j)_F, \text{Finish}(j-1)_F) + \text{Length}(j)_F$

Fair Queuing (2)

- Suppose:

- Flow 1 and 3 use 1000B byte packets, flow 2 uses 300B packets
 - What will fair queuing do?



Fair Queuing (3)

- Suppose:
 - Flow 1 and 3 use 1000B packets, flow 2 uses 300B packets
 - What will fair queuing do?

Let $\text{Finish}(0)_F = 0$, queues backlogged $[\text{Arrive}(j)_F < \text{Finish}(j-1)_F]$


→ $\text{Finish}(1)_{F1} = 1000$, $\text{Finish}(2)_{F1} = 2000$, ...

$\text{Finish}(1)_{F2} = 300$, $\text{Finish}(2)_{F2} = 600$, $\text{Finish}(3)_{F2} = 900$, 1200, 1500, ...

→ $\text{Finish}(1)_{F3} = 1000$, $\text{Finish}(2)_{F3} = 2000$, ...

Fair Queuing (4)

- Suppose:
 - Flow 1 and 3 use 1000B byte packets, flow 2 uses 300B packets
 - What will fair queuing do?



➤ WFQ (Weighted Fair Queuing)

- WFQ is a useful generalization of Fair Queuing:

- Assign a weight, Weight_F , to each flow
 - Higher weight gives more bandwidth, e.g., 2 is 2X bandwidth
- Change computation of $\text{Finish}(j)_F$ to factor in Weight_F


$\text{Arrive}(j)_F$ = arrival time of j-th packet of flow F

$\text{Length}(j)_F$ = length of j-th packet of flow F

$\text{Finish}(j)_F = \max (\text{Arrive}(j)_F, \text{Finish}(j-1)_F) + \text{Length}(j)_F / \text{Weight}_F$

WFQ Example

- An example you can work through ...




Packet	Arrival time	Length	Finish time	Output order
A	0	8	8	1
B	5	6	11	3
C	5	10	10	2
D	8	9	20	7
E	8	8	14	4
F	10	6	16	5
G	11	10	19	6
H	20	8	28	8

Handwritten notes: x2 next to rows C, D, and E; a bracket under the first three columns labeled 'input'; and two arrows pointing to the 'Finish time' and 'Output order' columns.

Using WFQ

- Lots of potential!
 - Can prioritize and protect flows
 - A powerful building block
- Not yet a complete solution
 - Need to determine flows (user? application? TCP connection?)
 - Difficult to implement at high speed for many concurrent flows
 - Need to assign weights to flows

END



© 2013 D. Wetherall

Slide material from: TANENBAUM, ANDREW S.; WETHERALL, DAVID J., COMPUTER NETWORKS, 5th Edition, © 2011.
Electronically reproduced by permission of Pearson Education, Inc., Upper Saddle River, New Jersey